

G54MDP

Mobile Device Programming

Mobile Phone Architecture

Brainstorming

- Pick a mobile device (perhaps one that you own)
- Think about what the characteristics and specifications of that device are
 - CPU speed, display type, size, shape, network connectivity, battery life etc
 - Operating system, mode of interaction / UI
- How might these affect how we program for it?

Learning Outcomes

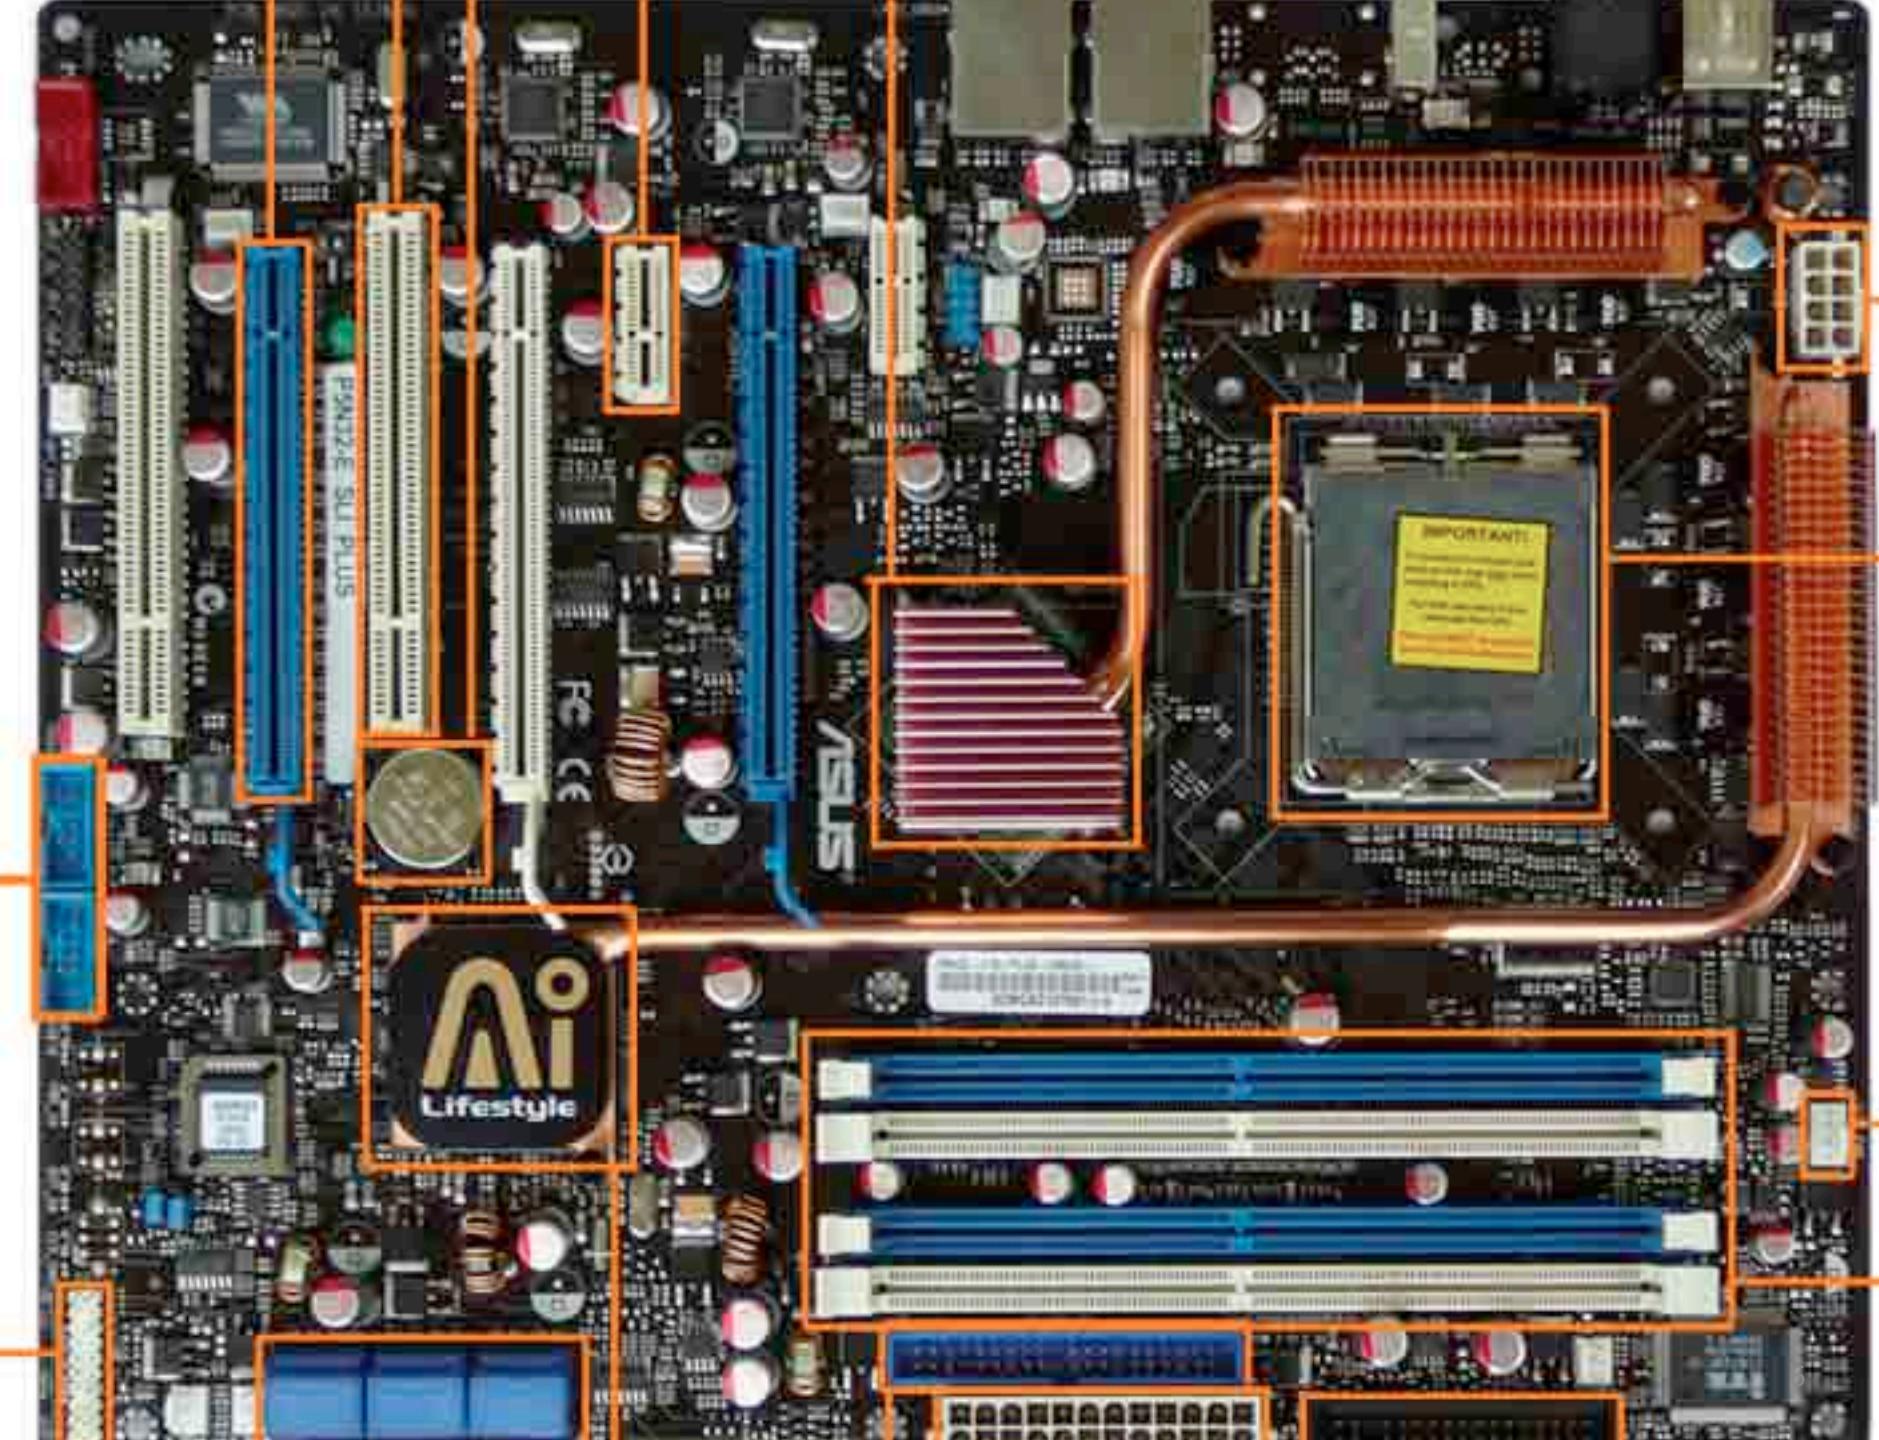
- By the end of this lecture you should...
- Understand the architectural differences between a PC and a mobile phone
- Have an overview of the main components of a mobile phone
- Understand some of the differences between ARM and x86 as relevant to mobile phones

Mobile Device Characteristics

- CPU ~1Gz
- GPU
- Memory
 - RAM 128MB-1GB
 - Flash Storage 16-64GB, internal / external via SD card
- Communications
 - Telephony
 - WiFi
 - Bluetooth
 - NFC
- Screen
- Audio
- User input
- Battery

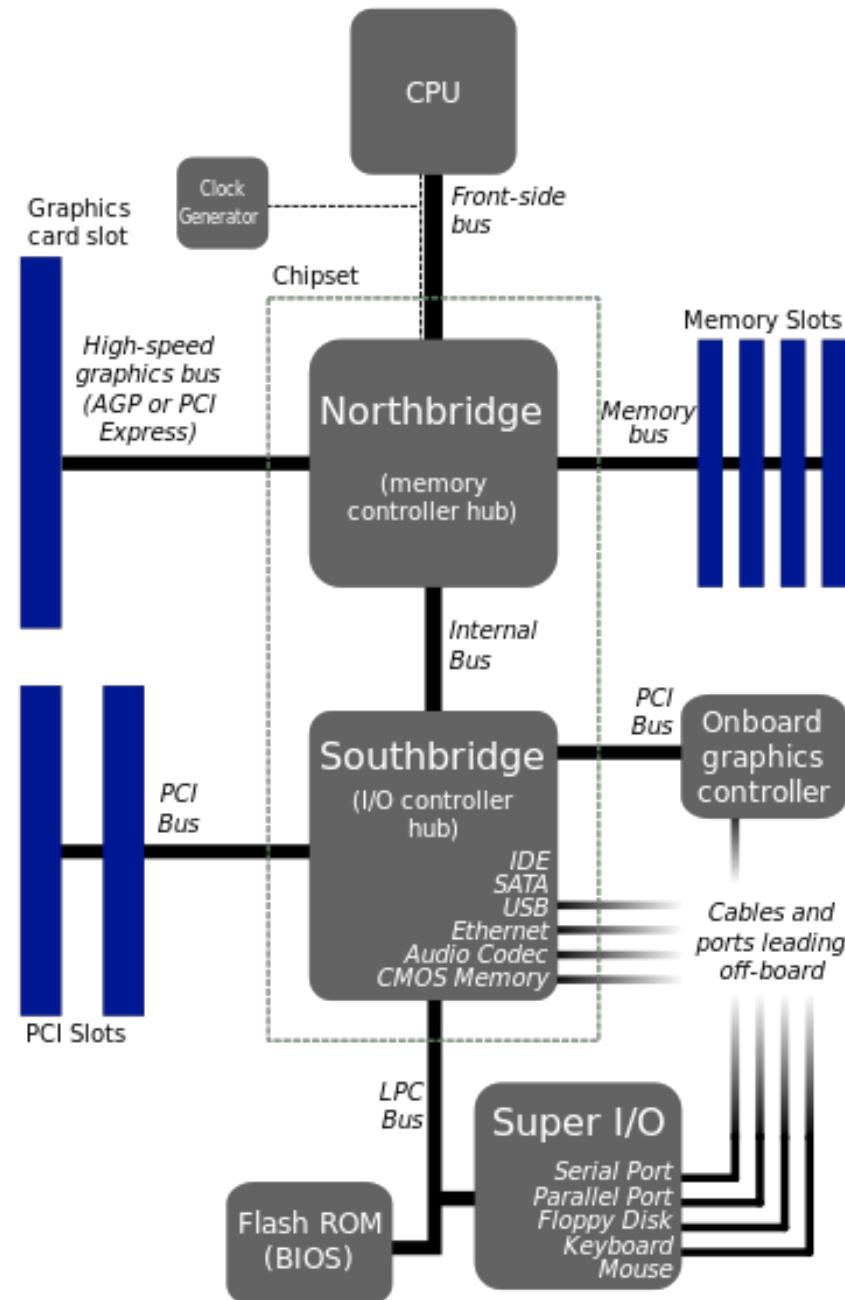
Sound familiar?

- In many ways, the technology inside a modern smart phone is conceptually very similar to that within a desktop PC
- It is the way that it is put together that is different



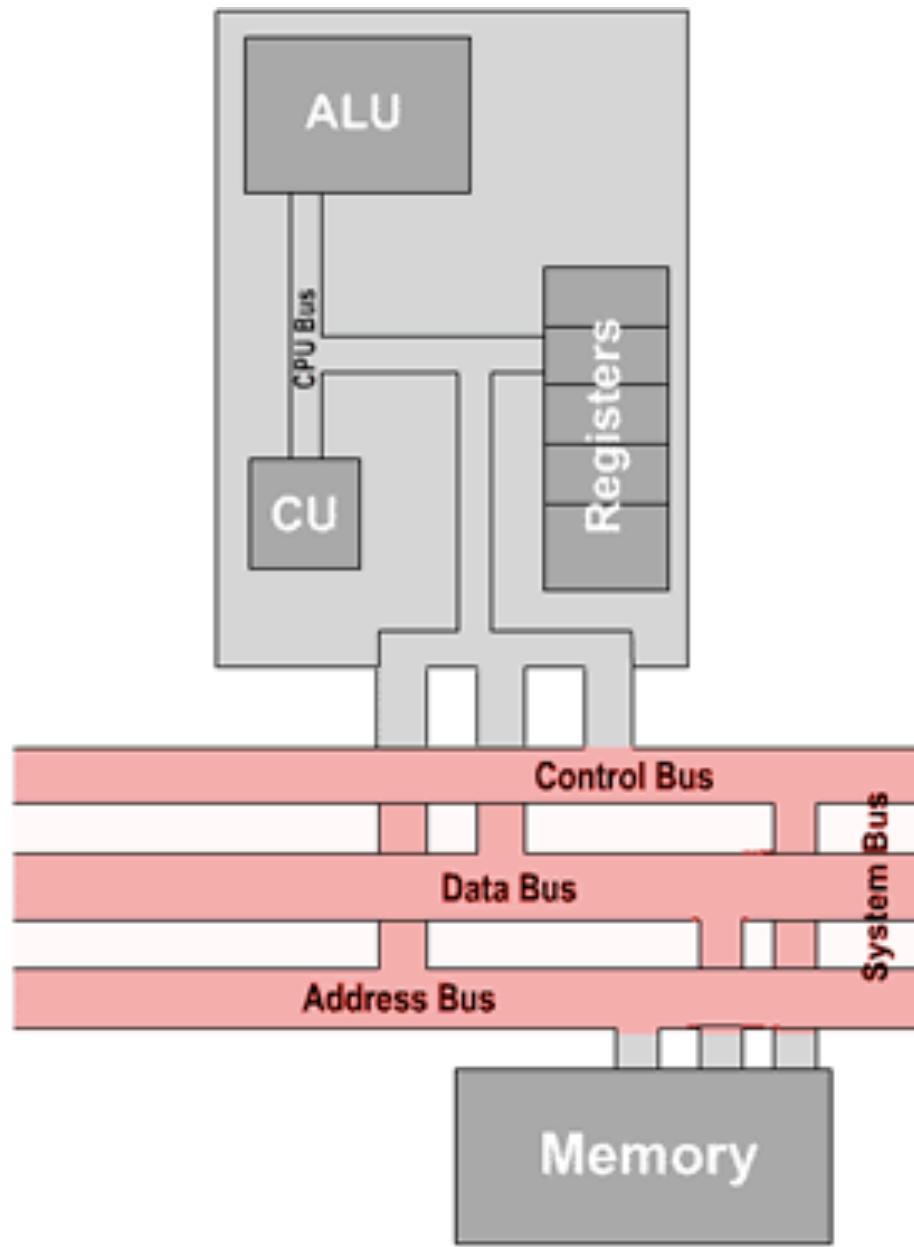
Inside a PC

- Several chips
- CPU
- Northbridge
- Southbridge
- RAM
- Several cards with extra bits
 - NIC, GPU, Audio
- Disks
 - HDD / SSD



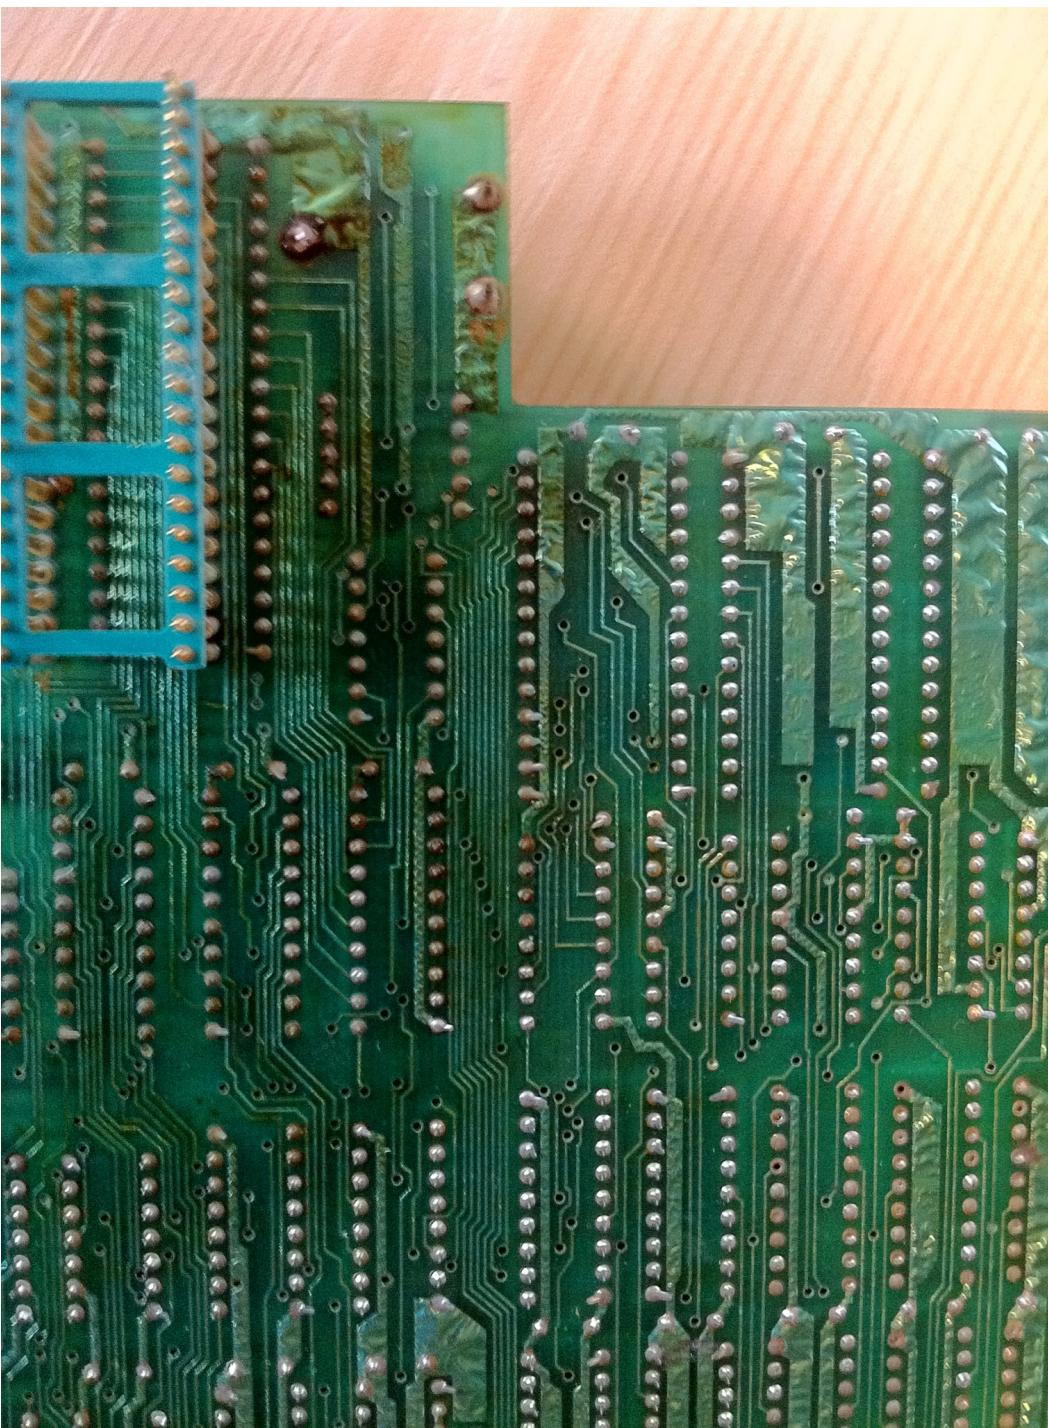
Connections

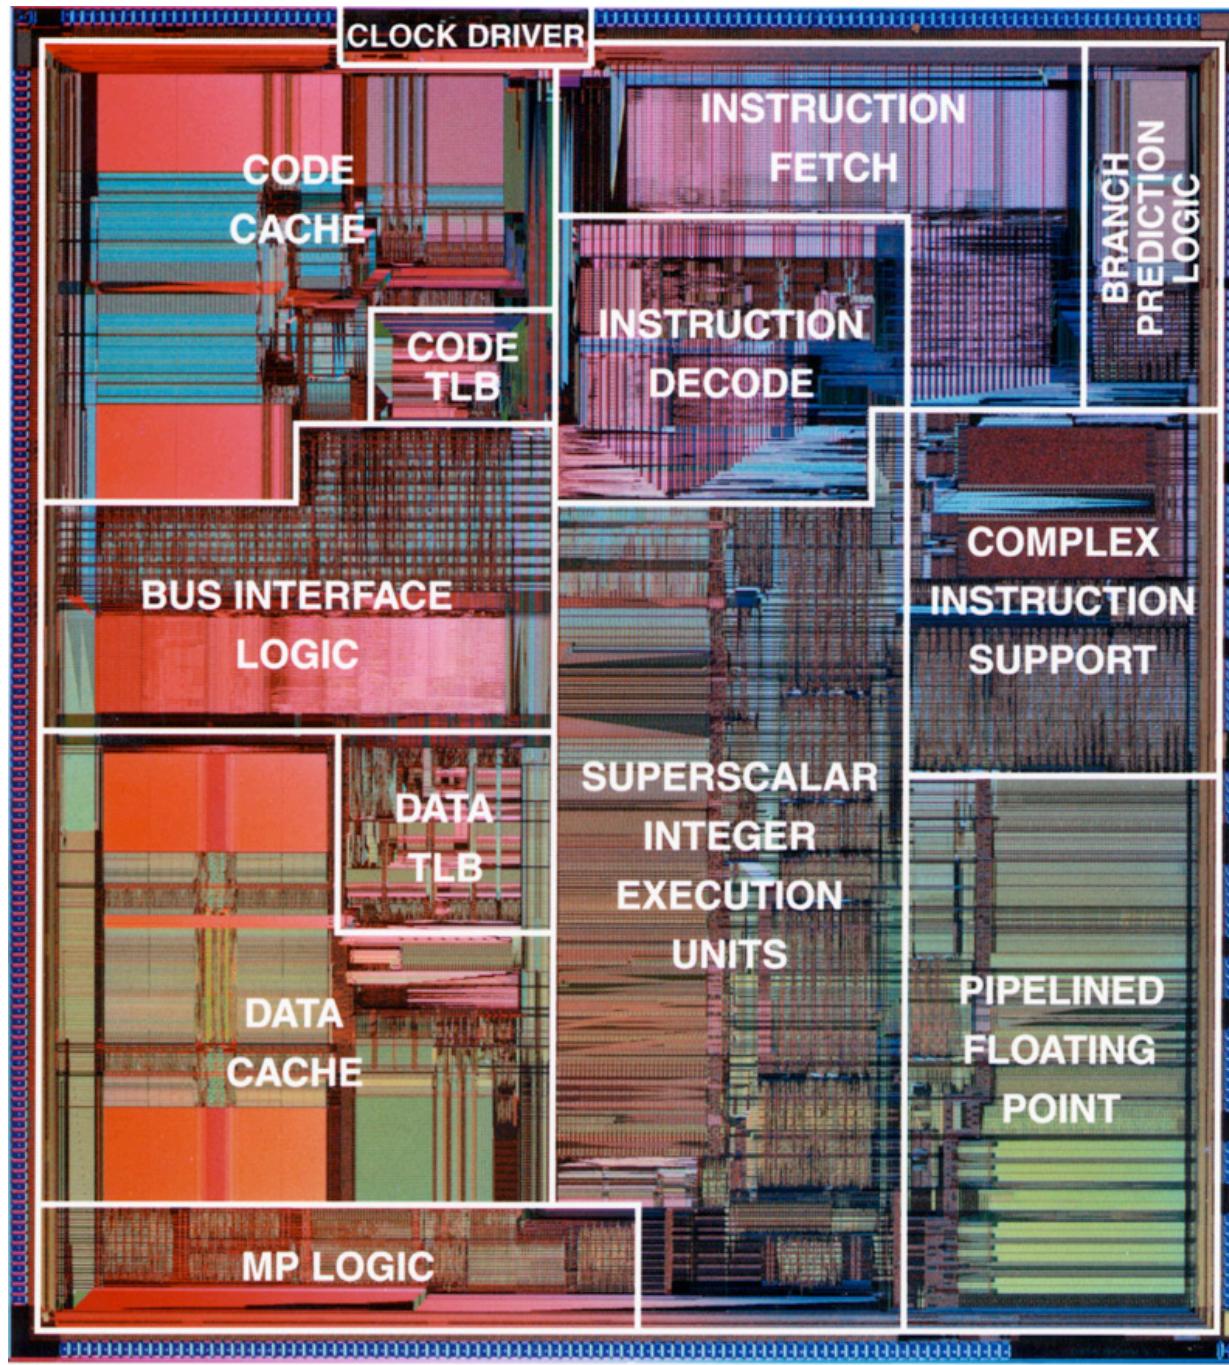
- CPU has a bus that connects it to the other devices / chips
- Originally, this would have been an address bus and a data bus, alongside some control buses
- RAM chips etc have similar connections
- Connect CPU to RAM / ROM
- Add some control logic and you have a PC



Connections

- Nowadays more advanced buses are used
 - E.g. HyperBus
- Bus width
 - The more lines (wires) the more data can be transferred
≈ faster bus
 - However, the number of pins used for the bus defines chip size
- Multiple chips and connections
 - Communications overhead / voltage drop = power
 - Clean architectural principle
 - Difficult to fit into a mobile device
 - Requires a different approach





Transistors

- A CPU is a collection of transistors
- All digital logic devices are built out of transistors
- 4 transistors will build a NAND gate
- Any logic circuit can be built using multiple NAND gates
- Getting better at getting lots of transistors on a chip
 - Moores law

System on a Chip

- Only use some of the transistors on a chip to form the CPU
- Use the rest to build the other required components of the system
- External pins connect directly to peripheral hardware, not core chips
- Called System on a Chip (SoC)
 - The whole system is literally on a single chip
 - Integrates several heterogeneous components on a chip
 - Aims to reduce communications overhead

System on a Chip

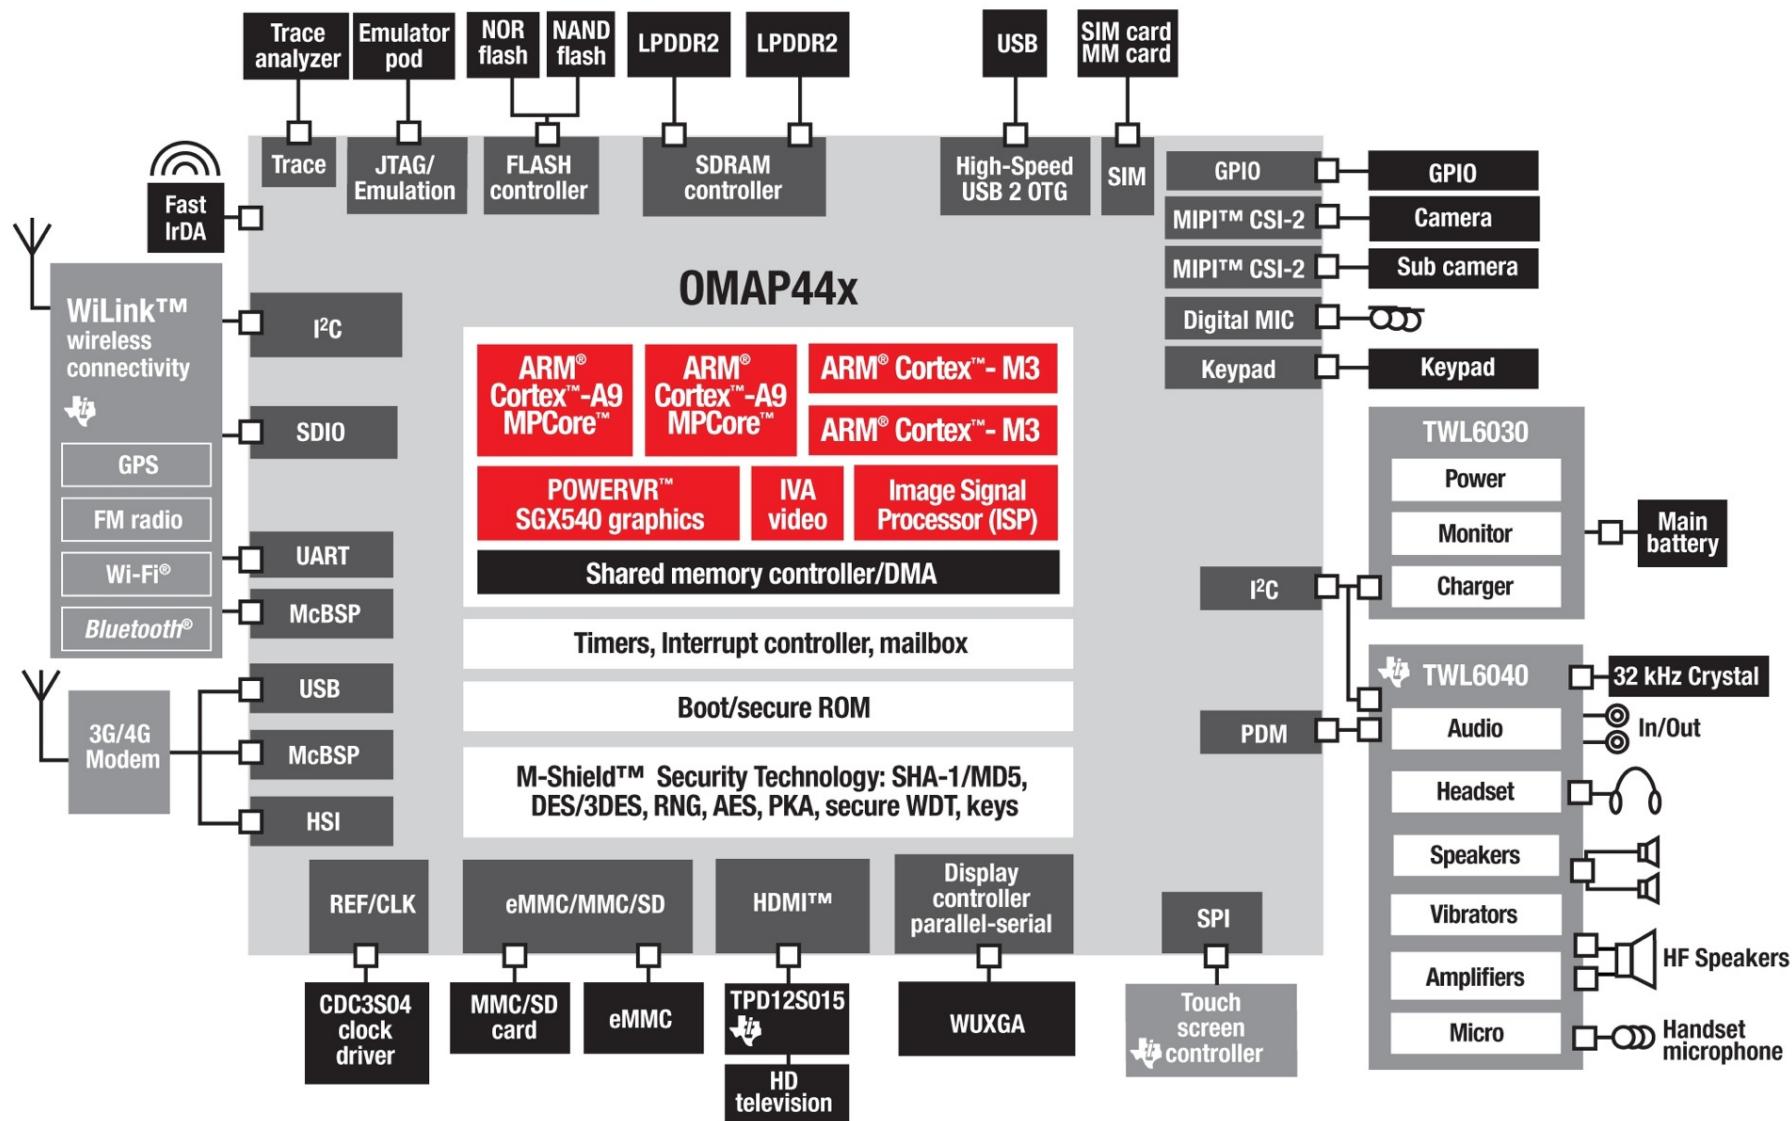
- Transistors provide computation, storage
 - Divide the chip into multiple communicating regions
- Built block by block from descriptions of separate parts
 - IP (intellectual property) cores
 - CPU core
 - GPU core
 - RAM core

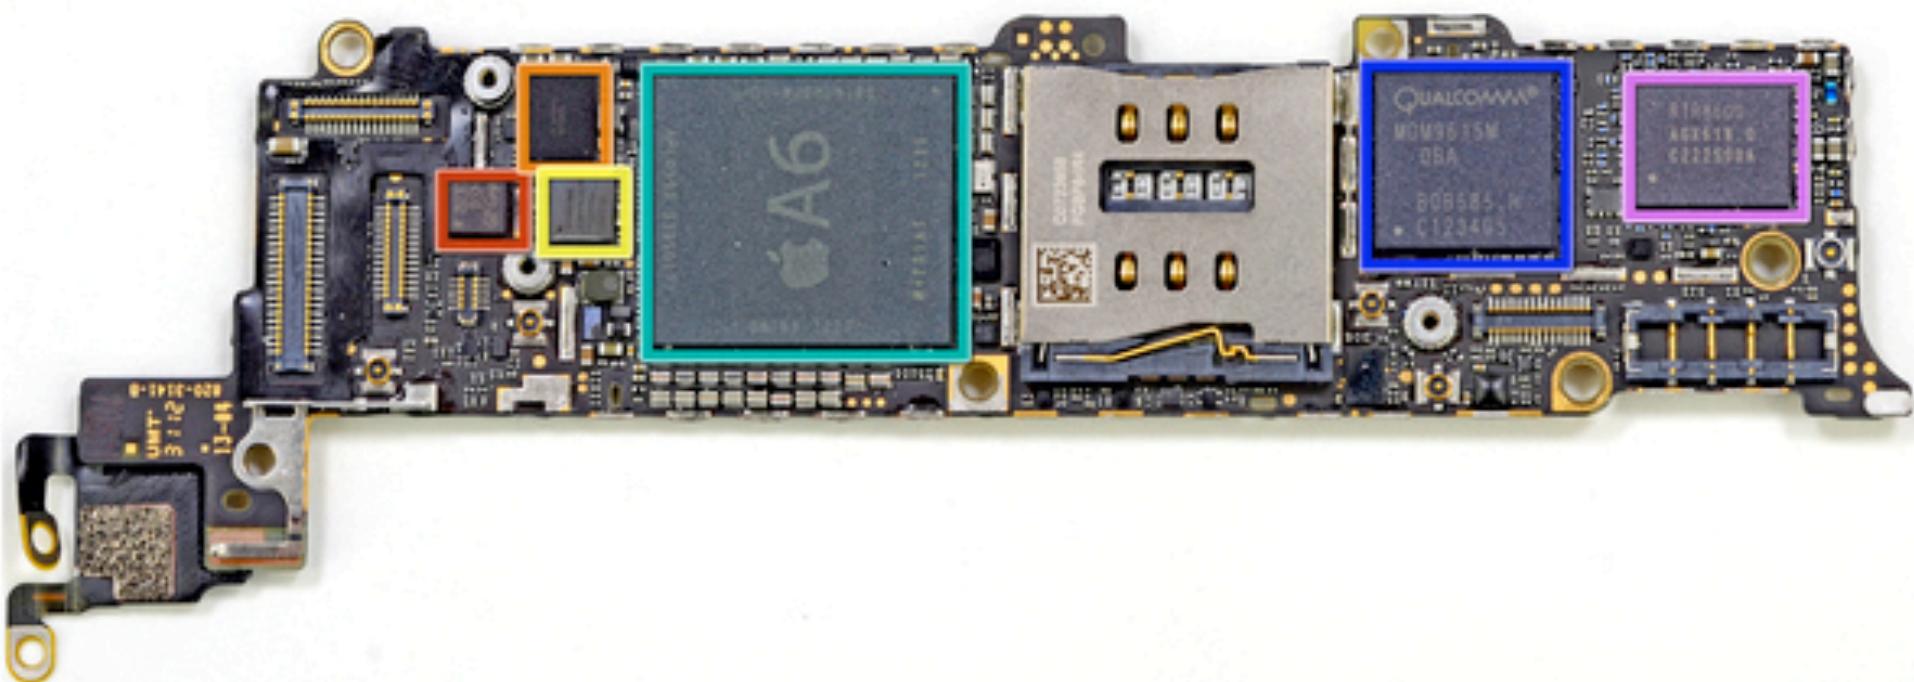
System on a Chip

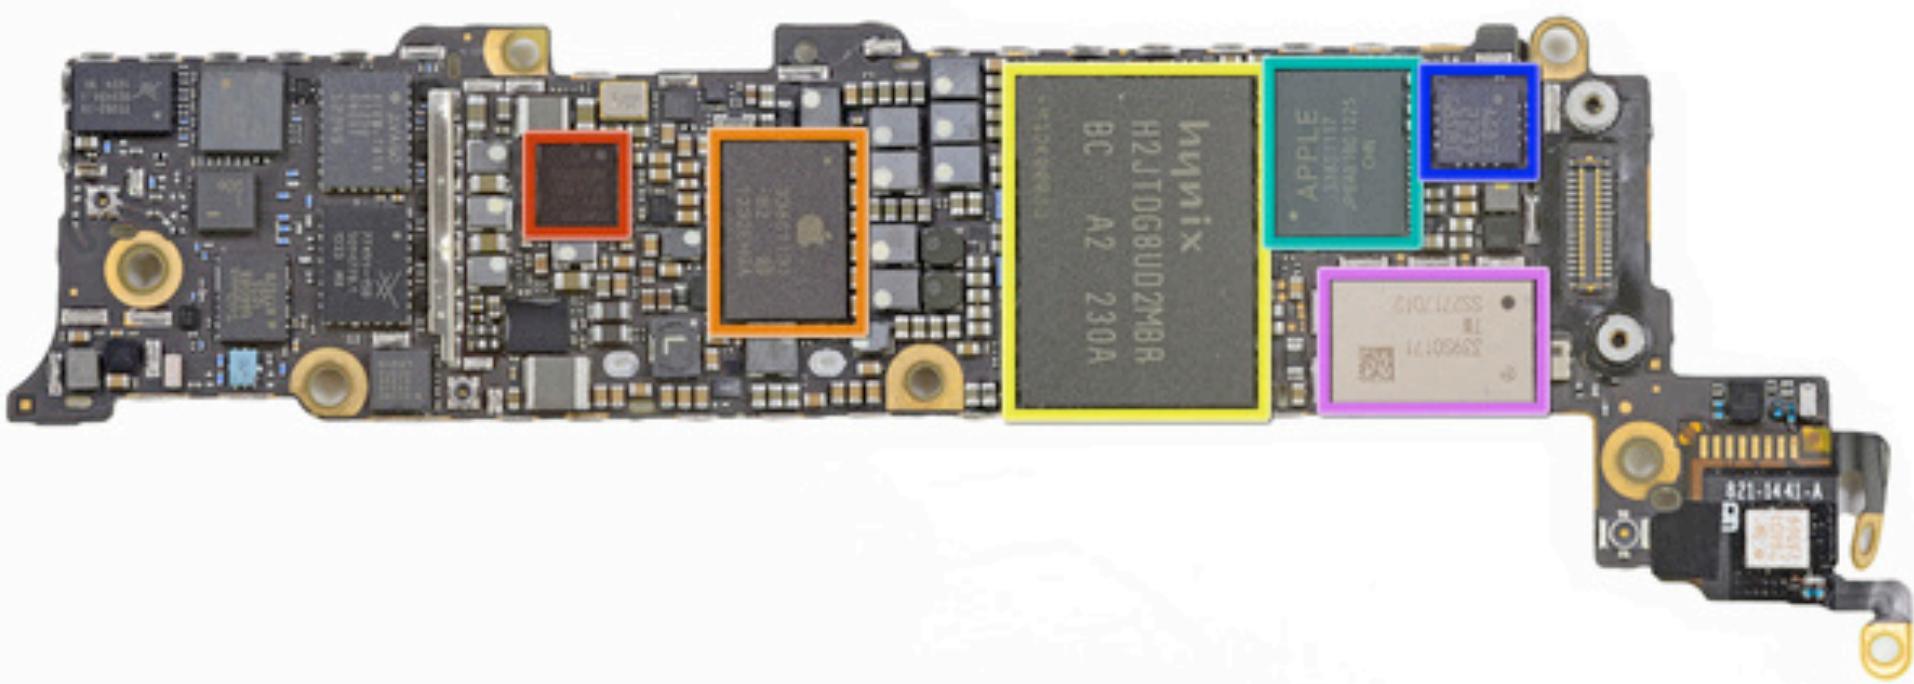
- Qualcomm Snapdragon
 - Nexus 4, Droid Razr, HTC Evo
- Apple A4/A5/A6/A7
 - iPhones
- Texas Instruments OMAP
 - Kindle Fire, Nook
- All share the same CPU block
 - ARM Cortex A8
- But may have different companion blocks

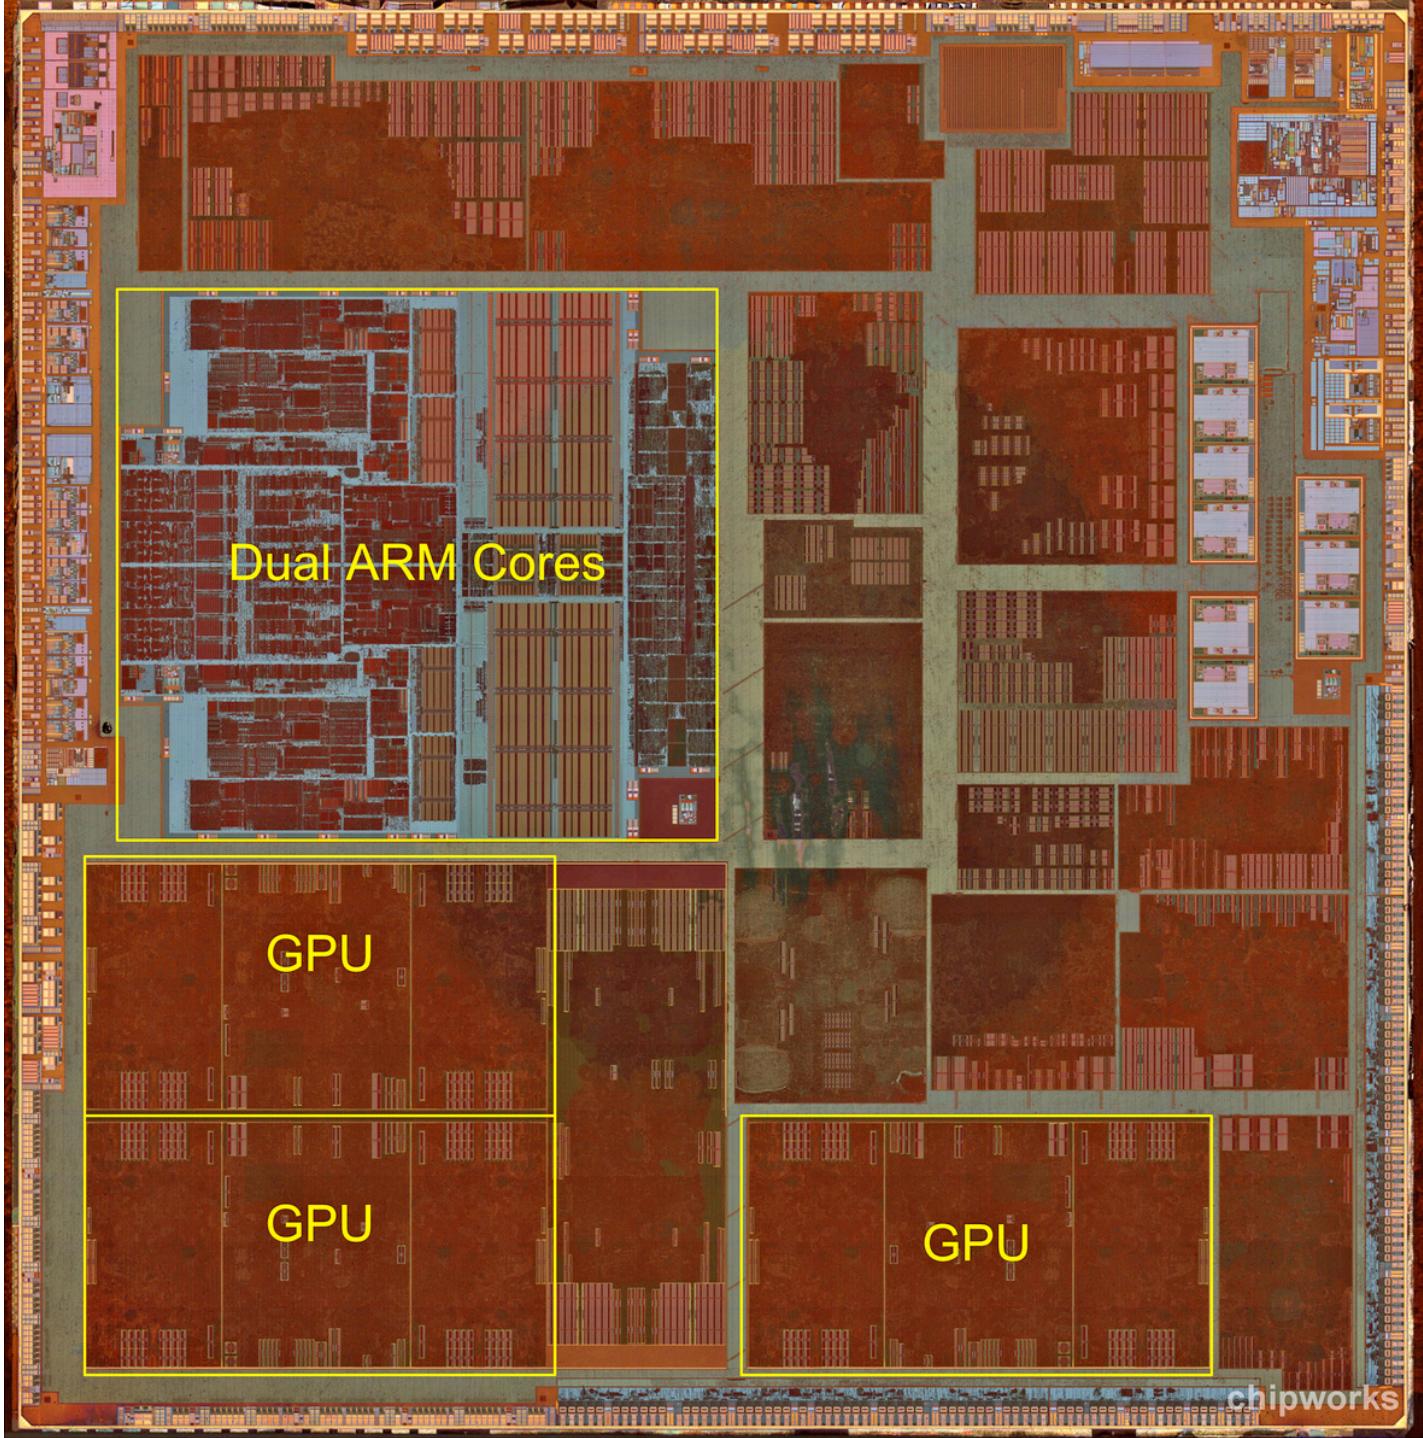
OMAP 44x0

- A typical SoC chip used in mobile phones
- Motorola Droid Bionic, Samsung Galaxy Nexus
- Contains two ARM Cortex-A9 CPU
- 3D GPU (PowerVR SGX)
- IVA Accelerator
- Image Signal Processor
- Some analogue components have a limited minimum size
 - Analogue vs digital – why?



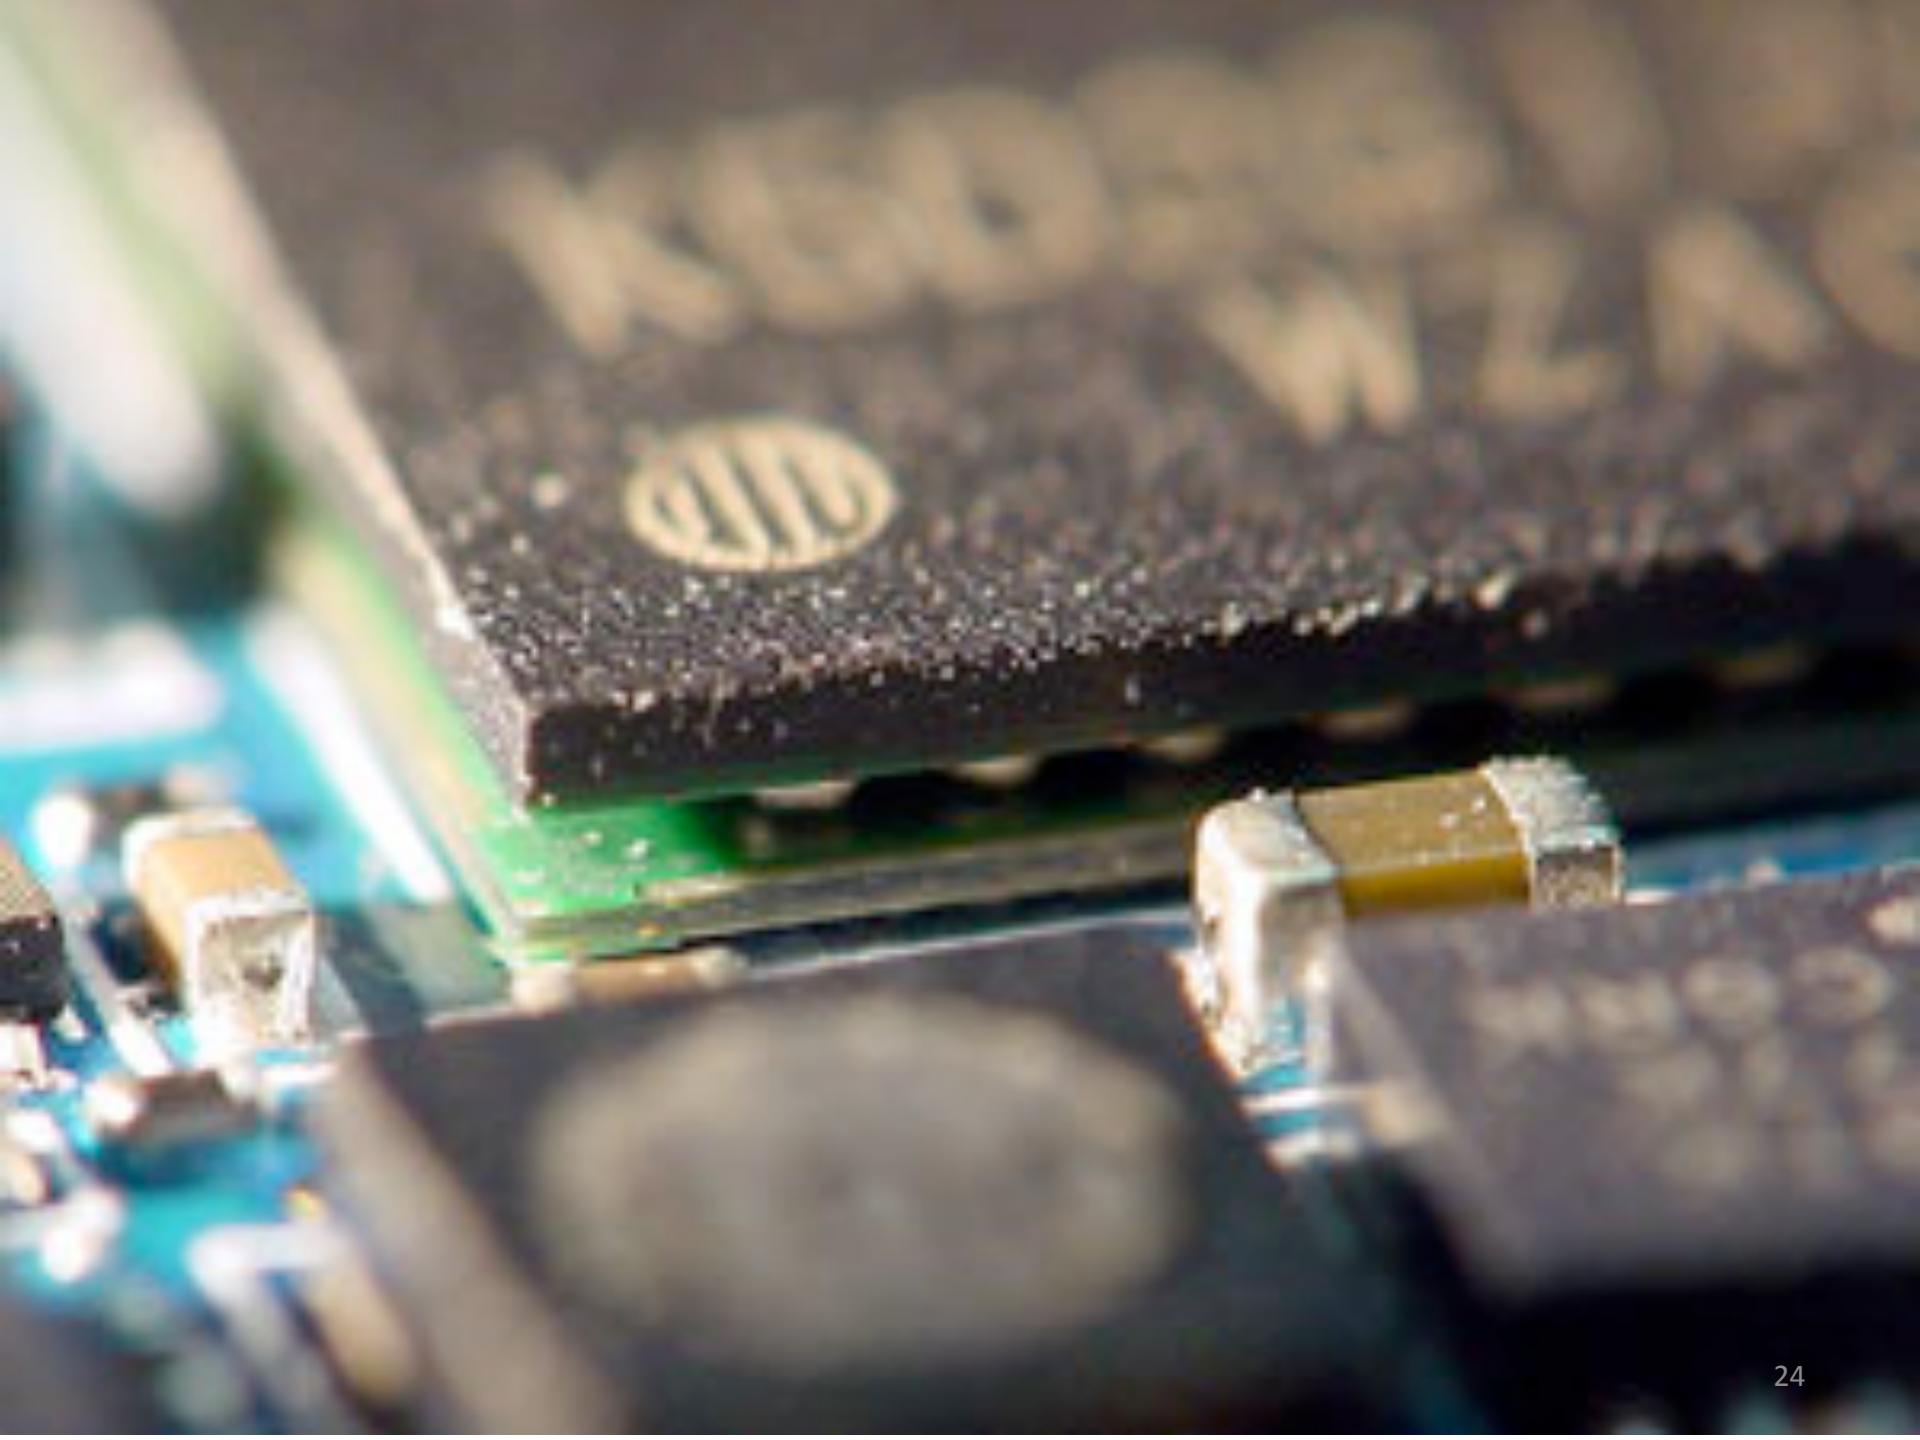


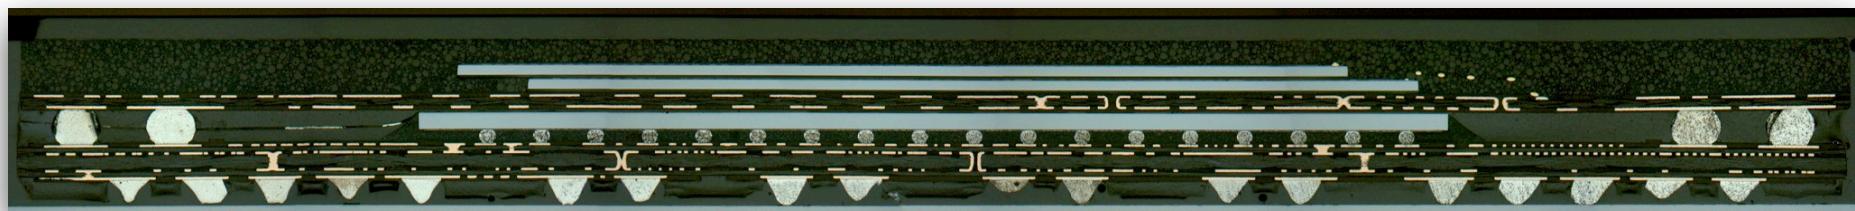




Package-on-Package

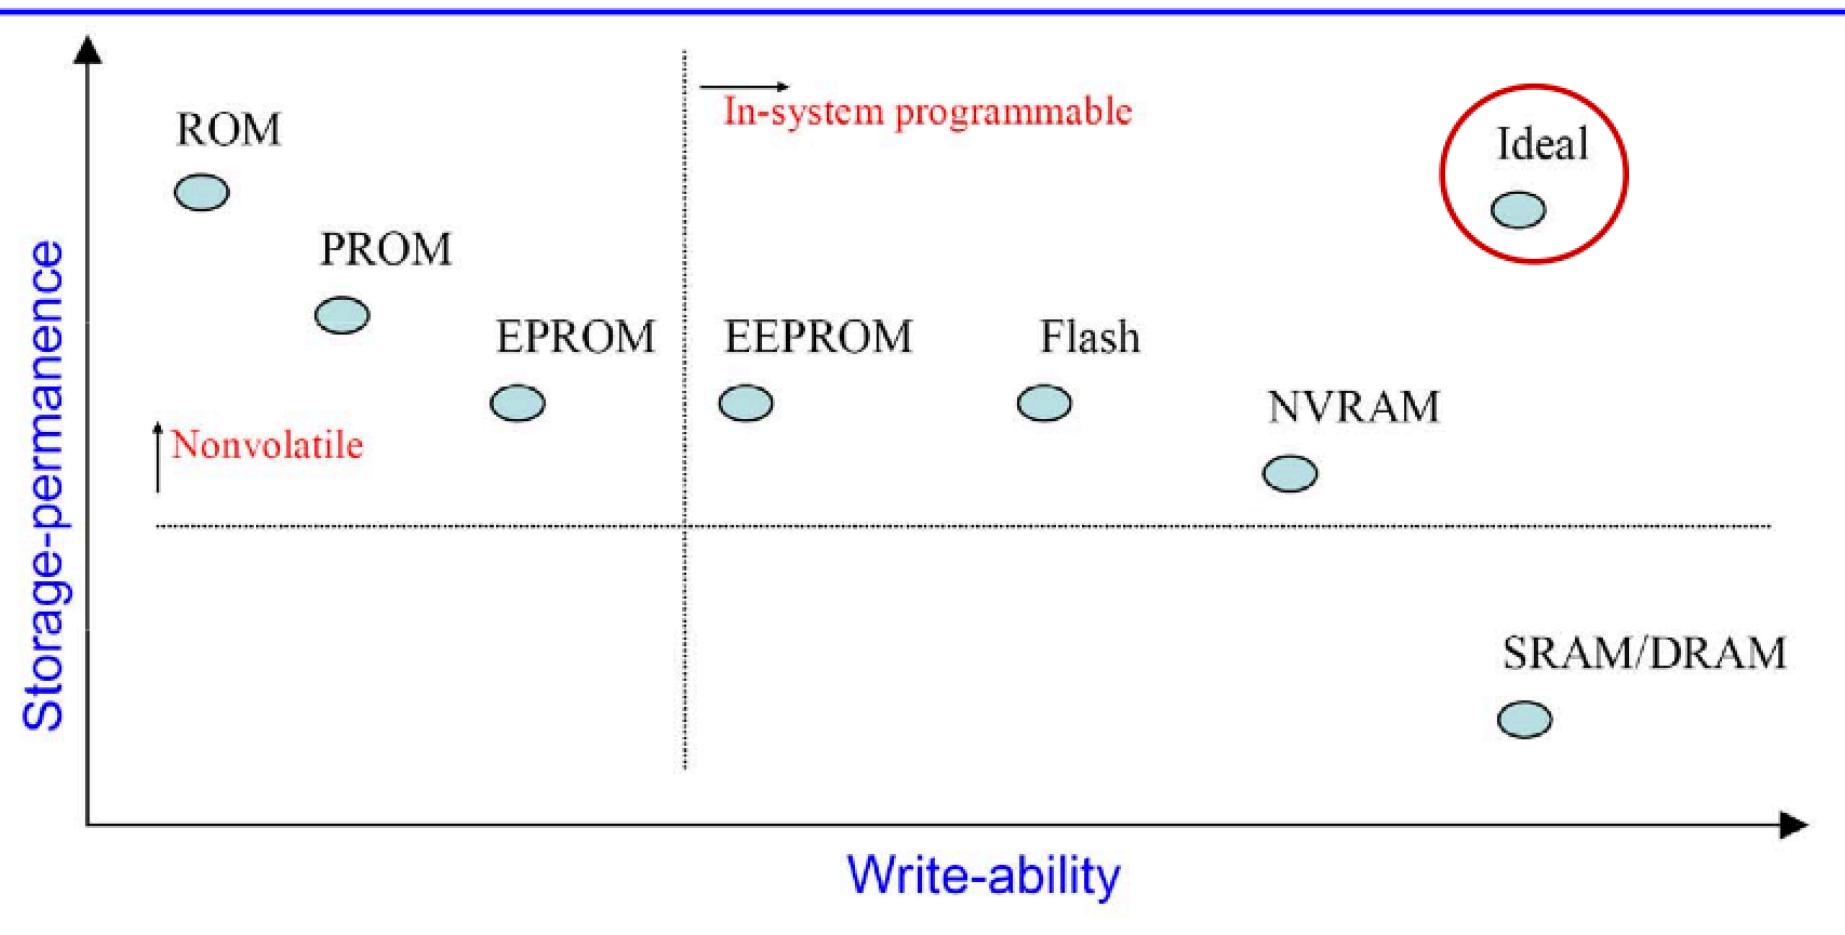
- What about RAM?
- Could add it to the SoC...
- But RAM uses a lot of space...
- Space we want to use for ‘useful’ stuff
- Separate package on top of the SoC
- Package-on-Package (or PoP)
 - Separate memory and logic production
 - “Stack” packages to minimise space





RAM

- RAM in a mobile device tends to be shared between multiple components on the main chip
 - C.f. laptop onboard GPU
 - Unlike separate motherboard RAM / GPU RAM
- Not all the RAM is available to the OS as “CPU RAM”
- No swap partition / page file
 - Why not?
 - What happens when you allocate all of the memory?
- Code assuming you don’t have very much



Mobile CPUs

- Almost all Mobile Devices use ARM CPUs
- Originated from Acorn computers in the late-1980s
- Acorn RISC Machine
- Spun out from Acorn in early-1990s
- Now, Advanced RISC machine

ARM CPU

- 32-bit CPU (increasingly 64-bit)
 - Multiple instruction sets
 - Older phones made use of Jazelle DBX / embedded systems
 - Run Java bytecode directly, dedicated chips
- Sold as a design, not a physical device
 - Great for SoC usage
 - Chip manufacturers can add ARM CPU to bespoke SoC vs buying a chip from Intel
- Aims
 - Fast and efficient
 - Good for mobile applications
 - Best use of battery – more instructions = more battery use
 - High code density
 - Less moving stuff around, best use of space

ARM CPU

- RISC-design (Reduced Instruction Set Chip)
- Only includes simple instructions that can be executed in one clock cycle
- Remove instructions you might expect in x86
 - e.g. divide instructions
 - If you need to divide, you roll your own in software
 - Or rather, the compiler does it for you

ARM CPU

Register names

r0	r1	r2	r3	r4	r5	r6	r7	r8	r9	r10	r11	r12	r13	r14	r15
----	----	----	----	----	----	----	----	----	----	-----	-----	-----	-----	-----	-----

- 16 registers
- Load/Store architecture
- Each instruction is 32-bits long
- Makes decoding easy / efficient
 - Vs x86 variable instruction length
- But means loading a constant into a register can be tricky

Opcode	Type	Decoded Instruction
0xE3A00000	Data Processing	MOV R0, #0
0xEF000002	Software Interrupt	SWI 2
0xEAFFFFFC	Branch Instruction	B -16
0xE0810002	Data Processing	ADD R0, R1, R2

C CPU - main thread, module winmine

0100368A	• 3B05 34530001 CMP EAX,DWORD PTR DS:[1005334]	
01003690	• 893D 64510001 MOV DWORD PTR DS:[1005164],EDI	
01003696	• ^75 0C JNZ SHORT winmine.010036A4	
01003698	• 3B00 38530001 CMP ECX,DWORD PTR DS:[1005338]	
0100369E	• ^75 04 JNZ SHORT winmine.010036A4	
010036A0	• 6A 04 PUSH 4	
010036A2	• ^EB 02 JMP SHORT winmine.010036A6	
010036A4	> 6A 06 PUSH 6	
010036A6	> 5B POP EBX	
010036A7	• A3 34530001 MOV DWORD PTR DS:[1005334],EAX	[1005334] = map-width
010036AC	• 890D 38530001 MOV DWORD PTR DS:[1005338],ECX	[1005338] = map-height
010036B2	• E8 1EF8FFFF CALL winmine.01002ED5	reset map memory
010036B7	• A1 A4560001 MOV EAX,DWORD PTR DS:[10056A4]	
010036BC	• 893D 60510001 MOV DWORD PTR DS:[1005160],EDI	
010036C2	• A3 30530001 MOV DWORD PTR DS:[1005330],EAX	[1005330] = number of mines
010036C7	> FF35 34530001 PUSH DWORD PTR DS:[1005334]	push map-width to the stack
010036CD	• E8 6E020000 CALL winmine.01003940	mine width = randomized width [0 - mapwidth-1]
010036D2	• FF35 38530001 PUSH DWORD PTR DS:[1005338]	push map-height to the stack
010036D8	• 8BF0 MOV ESI,EAX	
010036D9	• 46 INC ESI	
010036D9	• E8 60020000 CALL winmine.01003940	mine width = mine width + 1
010036E0	• 40 INC EAX	mine height = randomized height [0 - mapheight-1]
010036E1	• 8BC8 MOV ECX,EAX	mine height = mine height + 1
010036E3	• C1E1 05 SHL ECX,5	cell address = 0x1005340 + 32 * height + width
010036E6	F68431 405300 TEST BYTE PTR DS:[ECX+ESI+1005340],80	test if cell position is already a mine
010036EE	• ^75 D7 JNZ SHORT winmine.010036C7	if so, re-do this iteration
010036F0	• C1E0 05 SHL EAX,5	
010036F3	• 8D8430 405300 LEA EAX,DWORD PTR DS:[EAX+ESI+1005340]	
010036FA	• 8008 80 OR BYTE PTR DS:[EAX],80	set cell address of mine to mine (80)
010036FD	• FF00 30530001 DEC DWORD PTR DS:[1005330]	decrease the number of mines
01003703	• ^75 C2 JNZ SHORT winmine.010036C7	repeat if there are mines left
01003705	• 8B00 38530001 MOV ECX,DWORD PTR DS:[1005338]	
01003708	• 0FAF0D 345300 IMUL ECX,DWORD PTR DS:[1005334]	
01003712	• A1 A4560001 MOV EAX,DWORD PTR DS:[10056A4]	
01003717	• 2BC8 SUB ECX,EAX	
01003719	• 57 PUSH EDI	
0100371A	• 893D 9C570001 MOV DWORD PTR DS:[100579C],EDI	
01003720	• A3 30530001 MOV DWORD PTR DS:[1005330],EAX	
01003725	• A3 94510001 MOV DWORD PTR DS:[1005194],EAX	
0100372A	• 893D A4570001 MOV DWORD PTR DS:[10057A4],EDI	
01003730	• 890D A0570001 MOV DWORD PTR DS:[10057A0],ECX	
01003736	• C705 00500001 MOV DWORD PTR DS:[1005000],1	
01003740	• E8 25FDFFFF CALL winmine.0100346A	
01003745	• 53 PUSH EBX	
01003746	• E8 05E2FFFFFF CALL winmine.01001950	Arg1 winmine.01001950
01003748	• 5F POP EDI	
0100374C	• 5E POP ESI	
0100374D	• 5B POP EBX	
0100374E	• C3 RETN	
0100374F	• 53 PUSH EBX	
	M01 CPU DWORD PTR DS:[FFCB+0]	

Constants

- Constants must fit into the 32-bit instruction width
- Can't therefore be the full 32-bits
- 8-bit + a 4-bit shift — gives a wide range of values
- Also a Move Negated instruction
- If that doesn't work, load from a literal pool

ARM Speed

- Does running at 1GHz mean it's going to be slow?
- GHz-speed tells us 'cycles per second'
 - Actual speed depends on how many cycles an instruction takes
 - ARM aims for one-cycle per instruction
- x86 instructions can take many cycles
 - If an ARM instruction takes 1 cycle at 1GHz
 - And an x86 instruction takes 3 cycles at 3GHz
 - Which is faster?
- Speed is not entirely down to clock speed
 - It's what you do with it

ARM Conditional Execution

- Often want to conditionally execute some code – while loops, for loops etc

```
cmp al,dl  
jg label1
```

- Traditional approach is to execute a compare
 - Branch if the condition is met
 - Branches are ‘expensive’ (difficult to make parallel efficiently)
- ARM allows any instruction to made conditional
 - Remove the need for an expensive branch
 - Smaller code footprint

ARM and Thumb

- Every ARM instruction is 32bits long
 - Can take up a lot of memory
 - Memory accesses take time
 - Can slow things down
- Thumb
 - 16-bit version of the ARM instruction set
 - Variable length instruction set
 - Took the most popular ARM instructions used and encoded them as 16-bit values
 - Why?

ARM and Thumb

- Speed
 - ARM instructions require 32bits to be read every instruction
 - Memory reads take time
 - Not all memory is 32bit wide
 - On smaller RAM width, takes multiple reads to get an instruction
- By making the instructions smaller gain a speed up
 - 8-bit memory two reads per instruction
 - 16-bit memory, one read per instruction
 - 32-bit memory, one read gives two instructions
- iPhone SDK generates Thumb by default

Floating Point

- Thumb doesn't encode FPU instructions
- CPU would have to branch to ARM instructions to execute it
 - This takes time
- FPU heavy code is better compiled to ARM, not Thumb
 - Assuming the device has an FPU!

ARM big.LITTLE

- ARM's latest processor contains two cores
 - 2.0GHz quad core optimised for performance (big core)
 - 1.0GHz quad core optimised for energy efficiency (LITTLE core)
- Two Cores are architecturally consistent
- System can switch between the two as appropriate for the task in hand

Next time

- Introduction to Android development tools
- Android as an operating system
 - What is it?